

COMPUTER ENGINEERING MASTER DEGREE

COMPUTER ARCHITECTURE

PROJECT DISCUSSION

INTEGER FACTORIZATION

PROFESSORS

ANTONIO COSIMO PRETE
ANTONIO DI TECCO

GROUP MEMBERS

GIOVANNI LIGATO
ALICE ORLANDINI



University of Pisa


SYLLABUS




ALGORITHM



GOALS




CPU



GPU

ALGORITHM



INTEGER FACTORIZATION

«Every positive **integer** can be written
uniquely as a product of **primes**»

MATHEMATICAL FORMULATION

TRIAL DIVISION

[SEQUENTIAL]

INPUT

N : Number to be Factorized



while $N \bmod i == 0$

$N = N / i$

Add i to **Primes**

if $N \neq 1$

Add N to **Primes**

OUTPUT

Primes : Product of primes

INTEGER FACTORIZATION

«Every positive **integer** can be written
uniquely as a product of **primes**»

MATHEMATICAL FORMULATION


TRIAL DIVISION

[PARALLEL]

INPUT

N : Number to be Factorized

K : Number of Threads



OUTPUT

Primes : Product of primes



Mutex

GOALS



GOALS




Execution Time

\leq **1** second



For numbers up to

18
DIGITS

e.g. **975734686214396237**




DEVELOPER



Develop a GPU version that
outperforms the CPU version

CPU

IMPLEMENTATION




11th Gen **intel**® Core(TM) **i5**-11400

x Cores per socket: **6**

= Threads per core: **2**

TOTAL # LOGICAL CORES: **12**



L1**d** Cache: **288** KiB

L1**i** Cache: **192** KiB

L2 Cache: **3** MiB

L3 Cache: **12** MiB

Base Frequency: **2.60** GHz

Max Turbo Frequency: **4.40** GHz

HARDWARE DETAILS

TOOLS




Execution Time

```
chrono::steady_clock::time_point start = chrono::steady_clock::now();  
parallelTrialDivision(NUMBER, NUM_THREADS);  
chrono::steady_clock::time_point end = chrono::steady_clock::now();
```

```
chrono::milliseconds duration = chrono::duration_cast<chrono::milliseconds>(end - start);
```

Wall-Clock Time




start



end



Intel **VTune** Profiler



Visual Studio Code

FIRST EXECUTION

#THREADS : 12

$$\underbrace{975734686214396237}_{18 \text{ DIGITS}} = 748609 * 1303396948493$$

EXECUTION TIME
568340
MILLISECONDS

≈

9
MINUTES

>>

1
SECOND



REQUIREMENT
NOT
SATISFIED

LET'S GO INTO
DETAIL!




EXECUTION TIME

FIRST VERSION

#ITERATIONS : 30

$599999999999991 = 3 * 199999999999997$
(14 DIGITS)




SPEED-UP

FIRST VERSION

#ITERATIONS : 30


599999999999991 = 3 * 199999999999997
(14 DIGITS)



(CI omitted due to non-normal distribution)

EFFECTIVE CPU UTILIZATION

Execution
Time [s]



Logical Cores simultaneously utilized

TOP HOTSPOTS

#THREADS : 12

0.588s

EXECUTION TIME

4.829s

CPU TIME

0s

PAUSED TIME

0s

(0.0% of CPU Time)

THREAD OVERSUBSCRIPTION


Function	CPU Time	% CPU Time
isPrime	4.808s	99.6%
findPrimesInRange	0.013s	0.3%
others	0.008s	0.2%

FUNCTION	CPU Time	% CPU Time
→ isPrime	4.808s	99.6%

```
bool isPrime(unsigned long long n) {  
  
    ...  
  
    for(unsigned long long i = 2; i * i <= n; ++i) {  
        if(n % i == 0) {  
            is_prime = false;  
            break;  
        }  
    }  
  
    ...  
}
```

CPU TIME

Total (%)	Self (s)
92,8%	4,483s



HOTSPOT

LET'S
OPTIMIZE!



OPTIMIZATIONS #1



MainThread: $N \bmod 2 == 0$

OPTIMIZATIONS #1



OPTIMIZATIONS #2

findPrimesInRange

```
for (unsigned long long i = start; i <= end; ++i) {  
    if (isPrime(i) && num % i == 0) {  
  
        int exponent = 0;  
        while (num % i == 0) {  
            exponent++;  
            num /= i;  
        }  
  
        {  
            lock_guard<mutex> lock(mtx);  
            primes.push_back({i, exponent});  
        }  
    }  
}
```

OPTIMIZATIONS #2

findPrimesInRange

```
for (unsigned long long i = start; i <= end; i += 2) {
    if (isPrime(i) && num % i == 0) {

        int exponent = 0;
        while (num % i == 0) {
            exponent++;
            num /= i;
        }

        {

            lock_guard<mutex> lock(mtx);
            primes.push_back({i, exponent});
        }
    }
}
```

OPTIMIZATIONS #2

findPrimesInRange

```
for (unsigned long long i = start; i <= end; i += 2) {  
    if (isPrime(i) && num % i == 0) {  
  
        int exponent = 0;  
        while (num % i == 0) {  
            exponent++;  
            num /= i;  
        }  
  
        {  
            lock_guard<mutex> lock(mtx);  
            primes.push_back({i, exponent});  
        }  
    }  
}
```

OPTIMIZATIONS #2

findPrimesInRange


```
for (unsigned long long i = start; i <= end; i += 2) {  
    if (num % i == 0) {  
  
        int exponent = 0;  
        while (num % i == 0) {  
            exponent++;  
            num /= i;  
        }  
  
        if (isPrime(i)) {  
            lock_guard<mutex> lock(mtx);  
            primes.push_back({i, exponent});  
        }  
    }  
}
```

EXECUTION TIME

OPTIMIZED VERSION

#ITERATIONS : 30

975734686214396237 = 748609 * 1303396948493
(18 DIGITS)



FIRST VERSION

EXECUTION TIME
568340
MILLISECONDS

≈

9
MINUTES

OPTIMIZED VERSION

MAX (OVER 30
ITERATIONS)
EXECUTION TIME
211
MILLISECONDS

<

1
SECOND

x2500
SPEED-UP

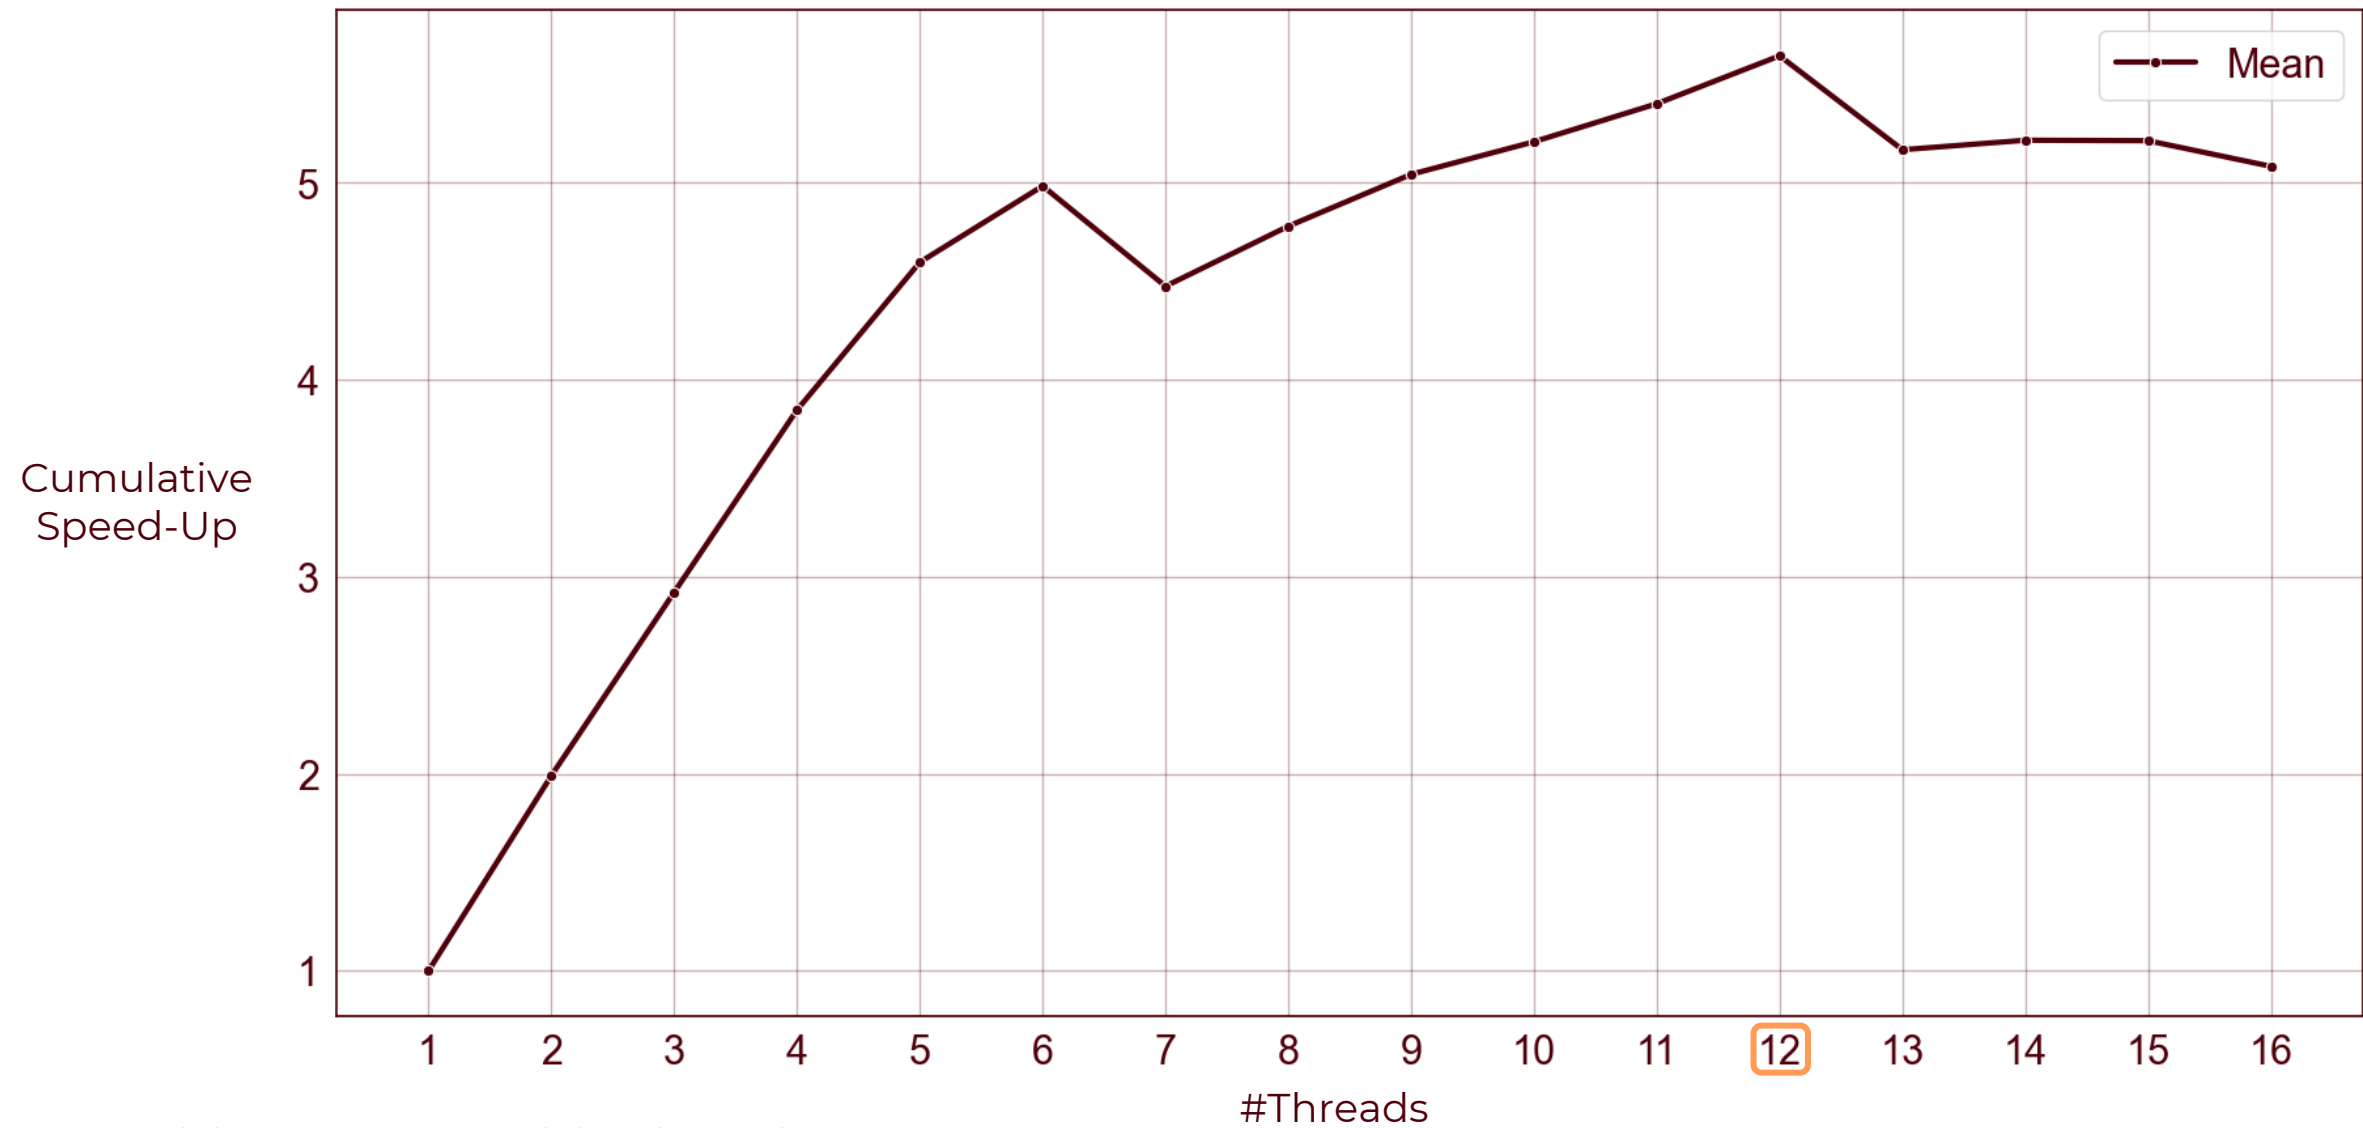
#THREADS : **12**

SPEED-UP

OPTIMIZED VERSION

#ITERATIONS : 30

975734686214396237 = 748609 * 1303396948493
(18 DIGITS)



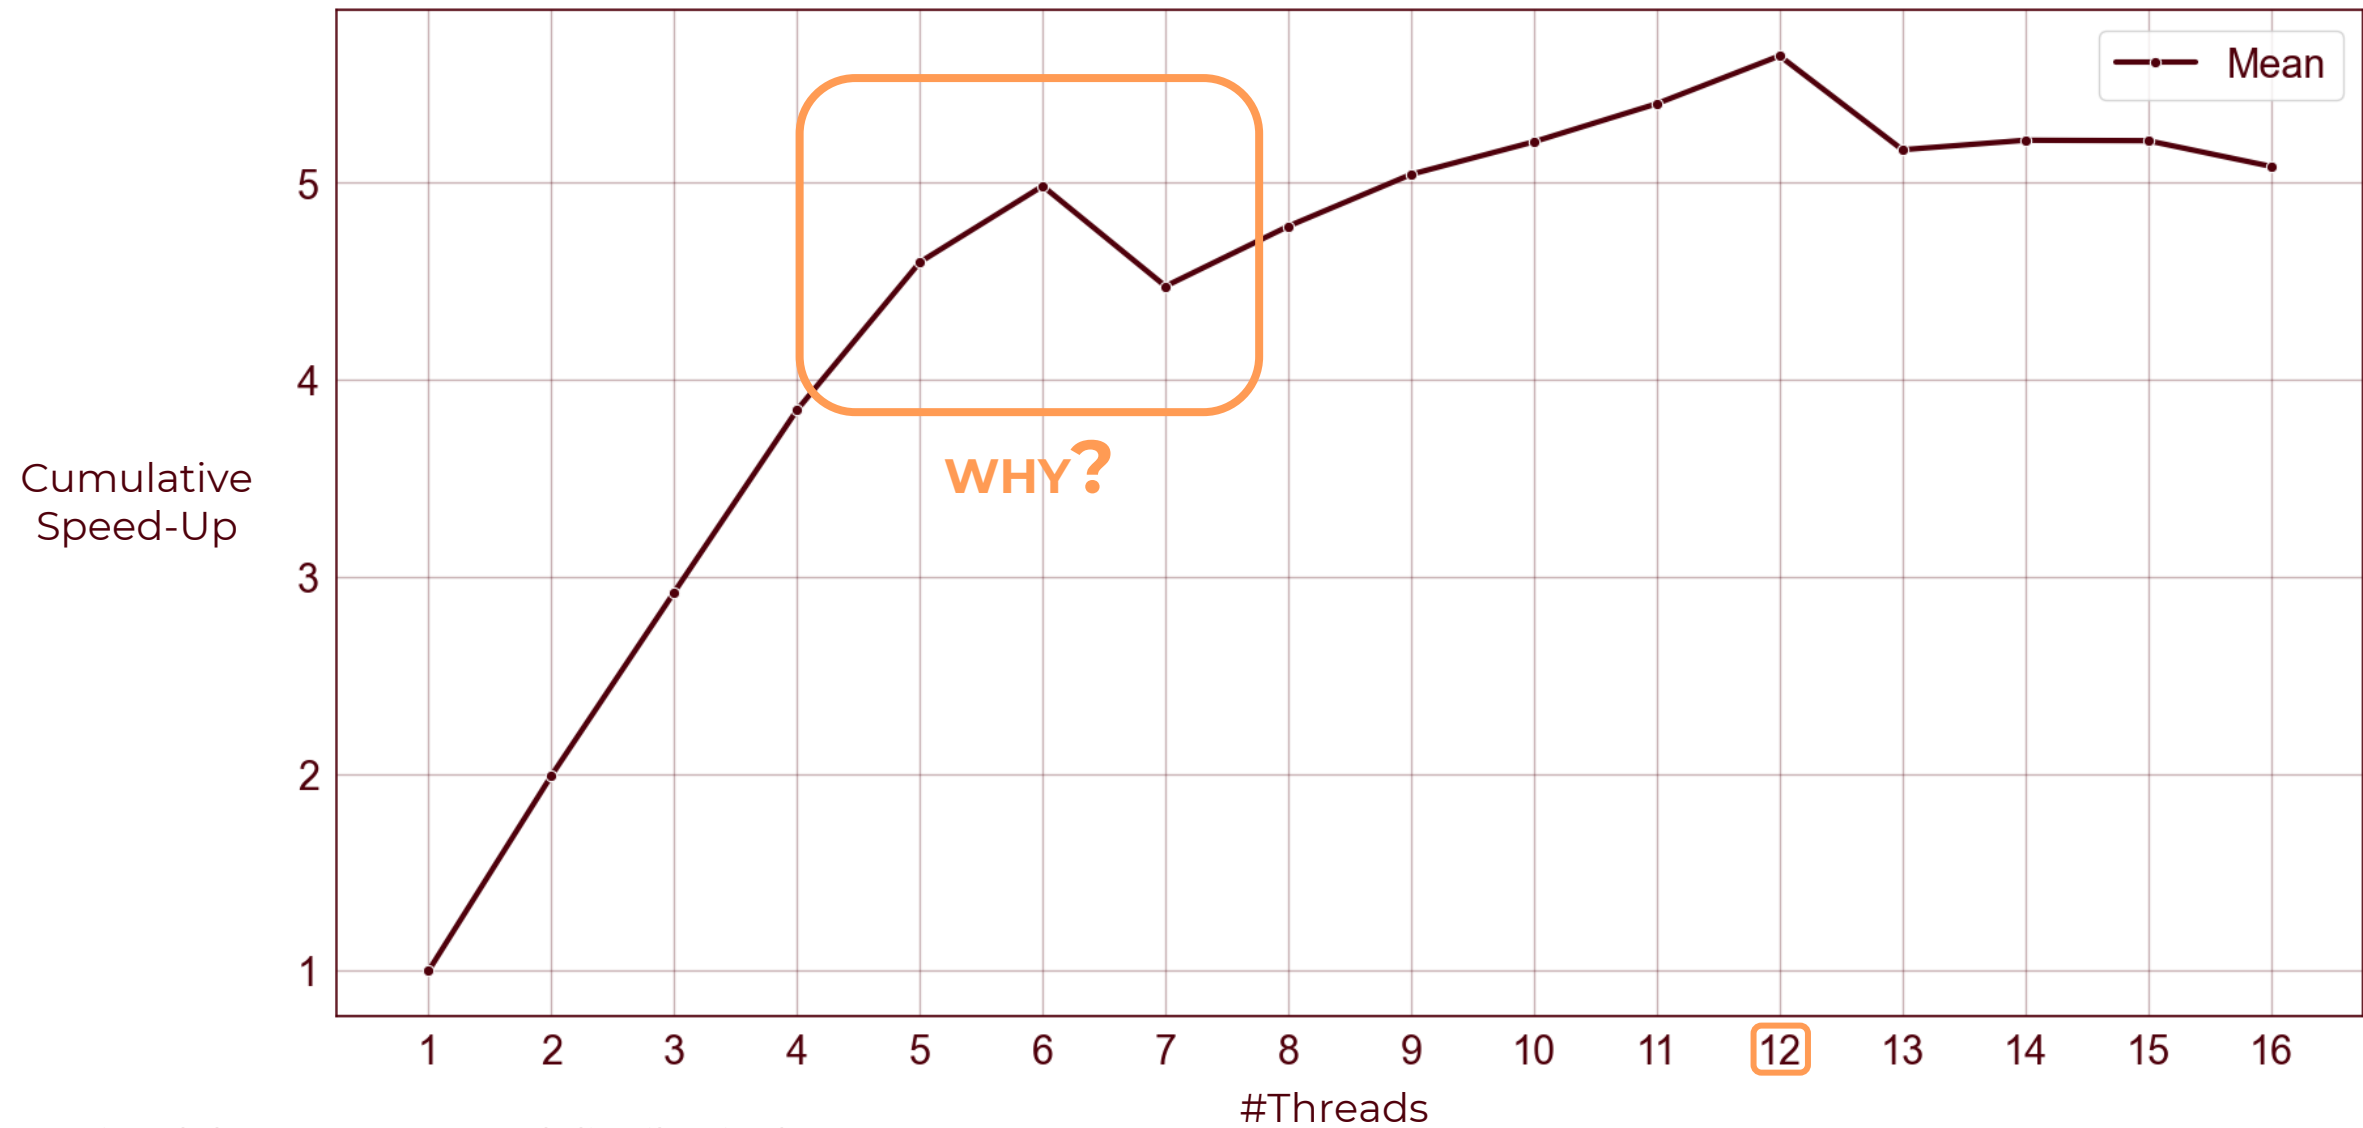
(CI omitted due to non-normal distribution)

SPEED-UP

OPTIMIZED VERSION

#ITERATIONS : 30

975734686214396237 = 748609 * 1303396948493
(18 DIGITS)



(CI omitted due to non-normal distribution)

SPEED-UP

OPTIMIZED VERSION

+

AFFINITY

#ITERATIONS : 30

975734686214396237 = 748609 * 1303396948493
(18 DIGITS)

CPU CORES

[2 THREADS x CORE]

#1

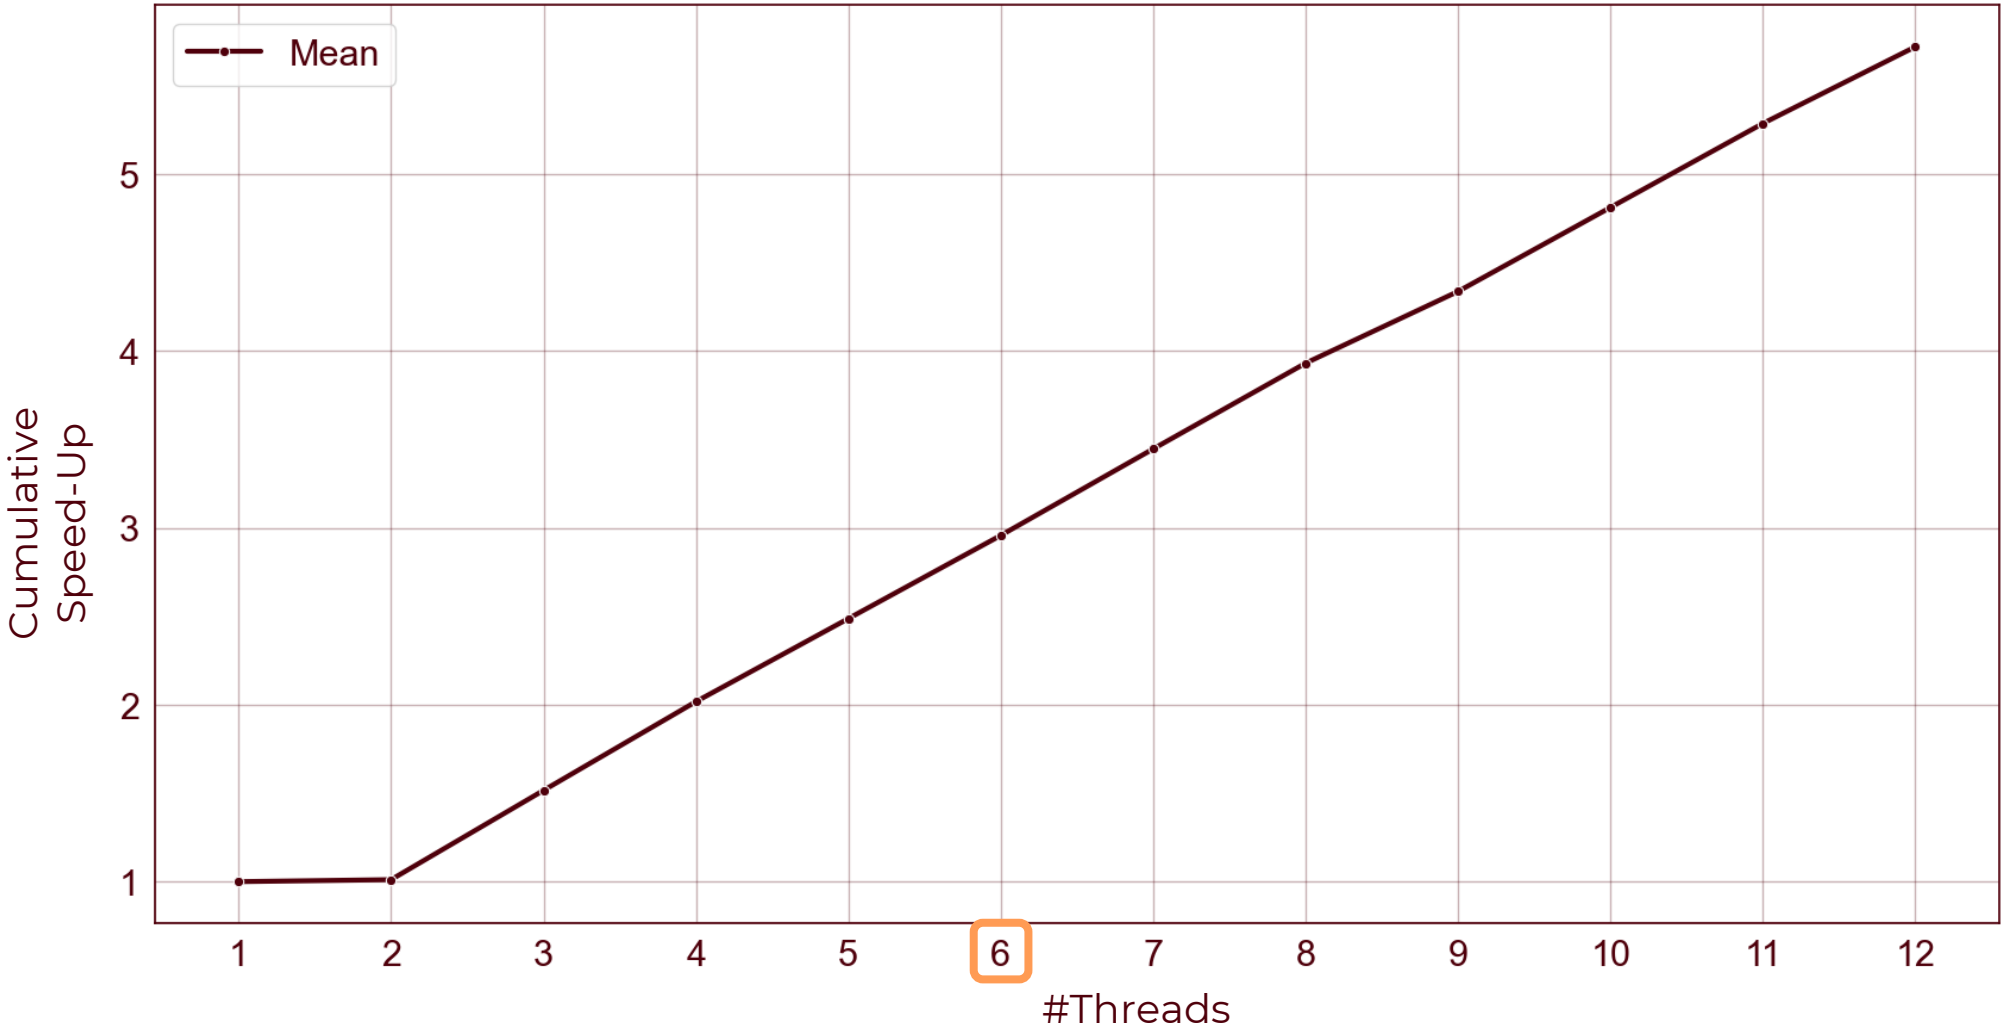
#2

#3

#4

#5

#6



(CI omitted due to non-normal distribution)


SPEED-UP

OPTIMIZED VERSION

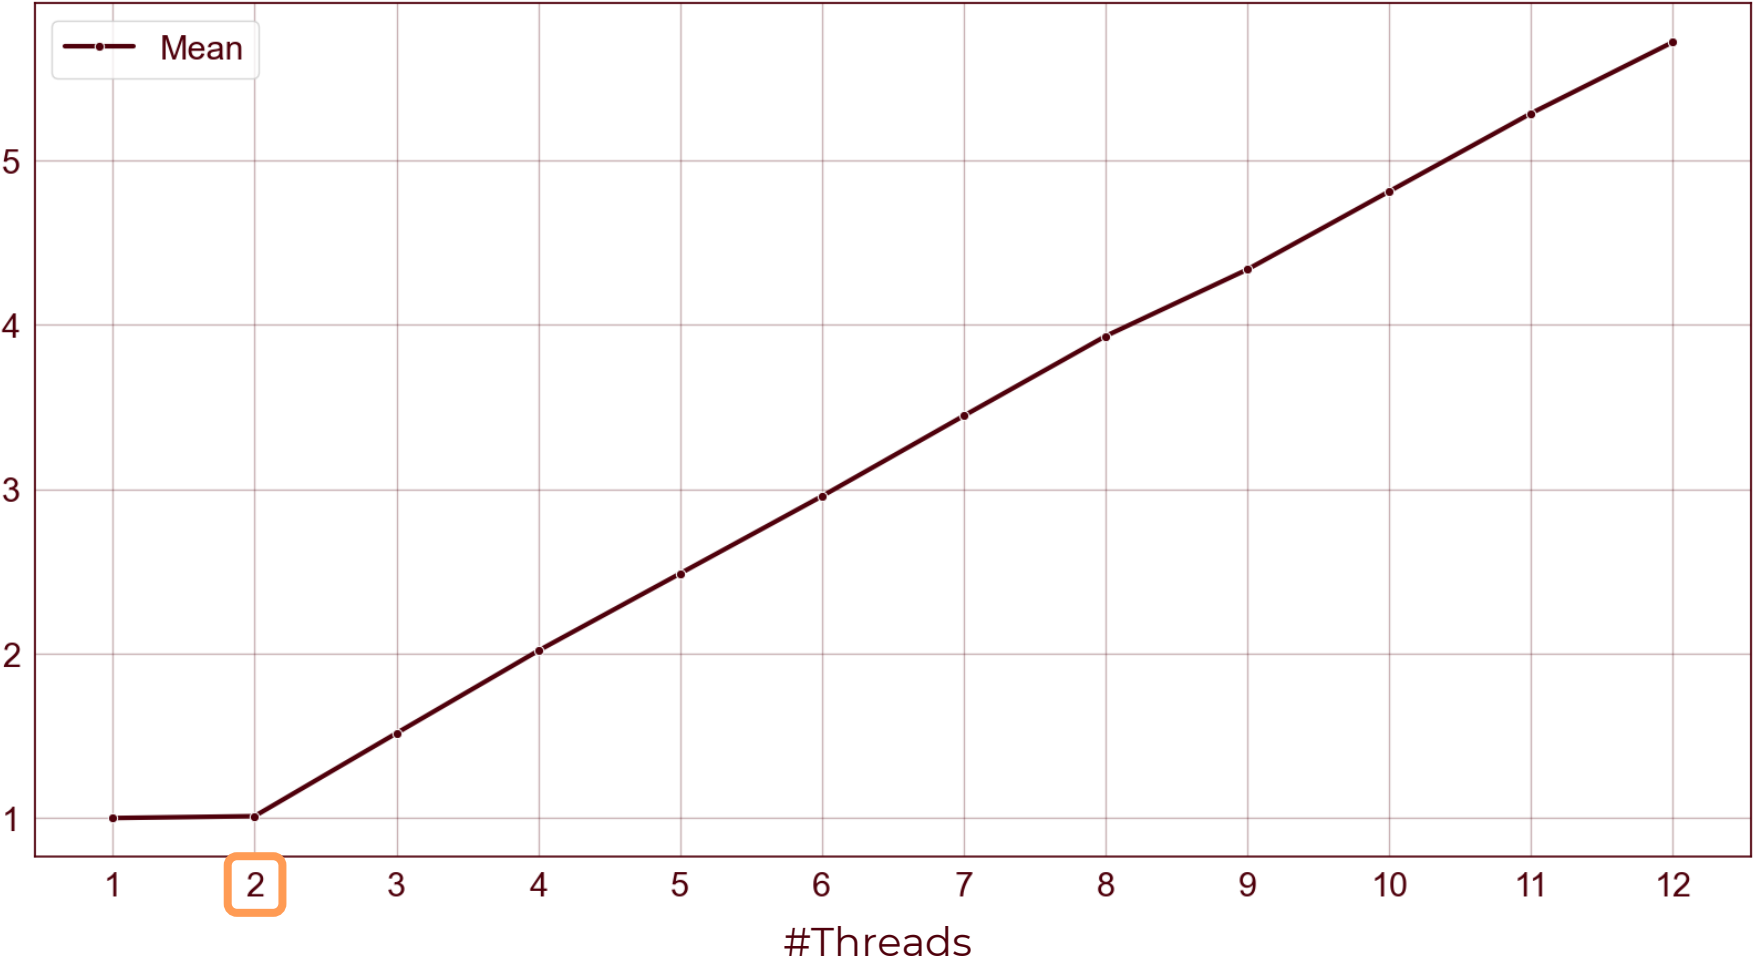
+

AFFINITY

CPU
CORES
[2 THREADS x CORE]




Cumulative
Speed-Up



THREAD 1

THREAD 2

`if (num % i == 0)`



DIVIDER
UTILIZATION
~ **56%**

SPEED-UP

OPTIMIZED VERSION

AFFINITY

#ITERATIONS : 30

975734686214396237 = 748609 * 1303396948493

(18 DIGITS)

CPU CORES

[2 THREADS x CORE]

#1


#2

#3

#4

#5

#6



#Threads

(CI omitted due to non-normal distribution)

SPEED-UP


OPTIMIZED VERSION

+

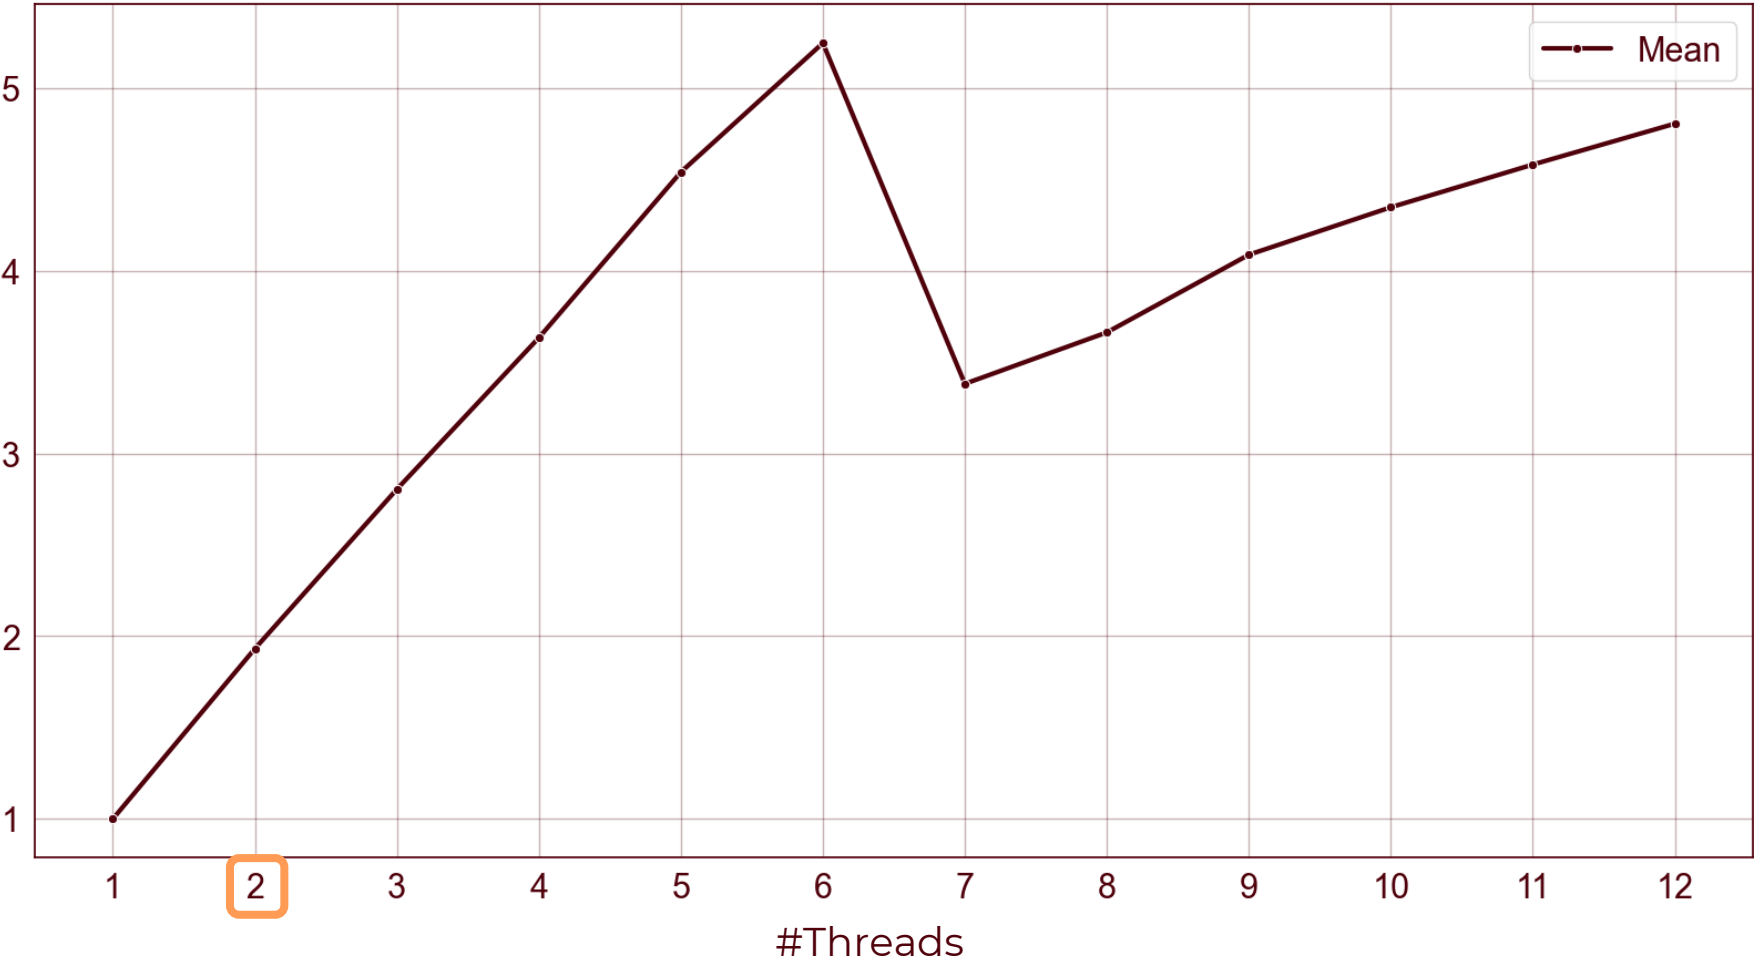
AFFINITY

CPU
CORES


[2 THREADS x CORE]




Cumulative
Speed-Up



THREAD 1 \longrightarrow `if (num % i == 0)` \longrightarrow



THREAD 2 \longrightarrow `if (num % i == 0)` \longrightarrow



DIVIDER
UTILIZATION
 \sim **92%**

EXECUTION TIME - SCALABILITY

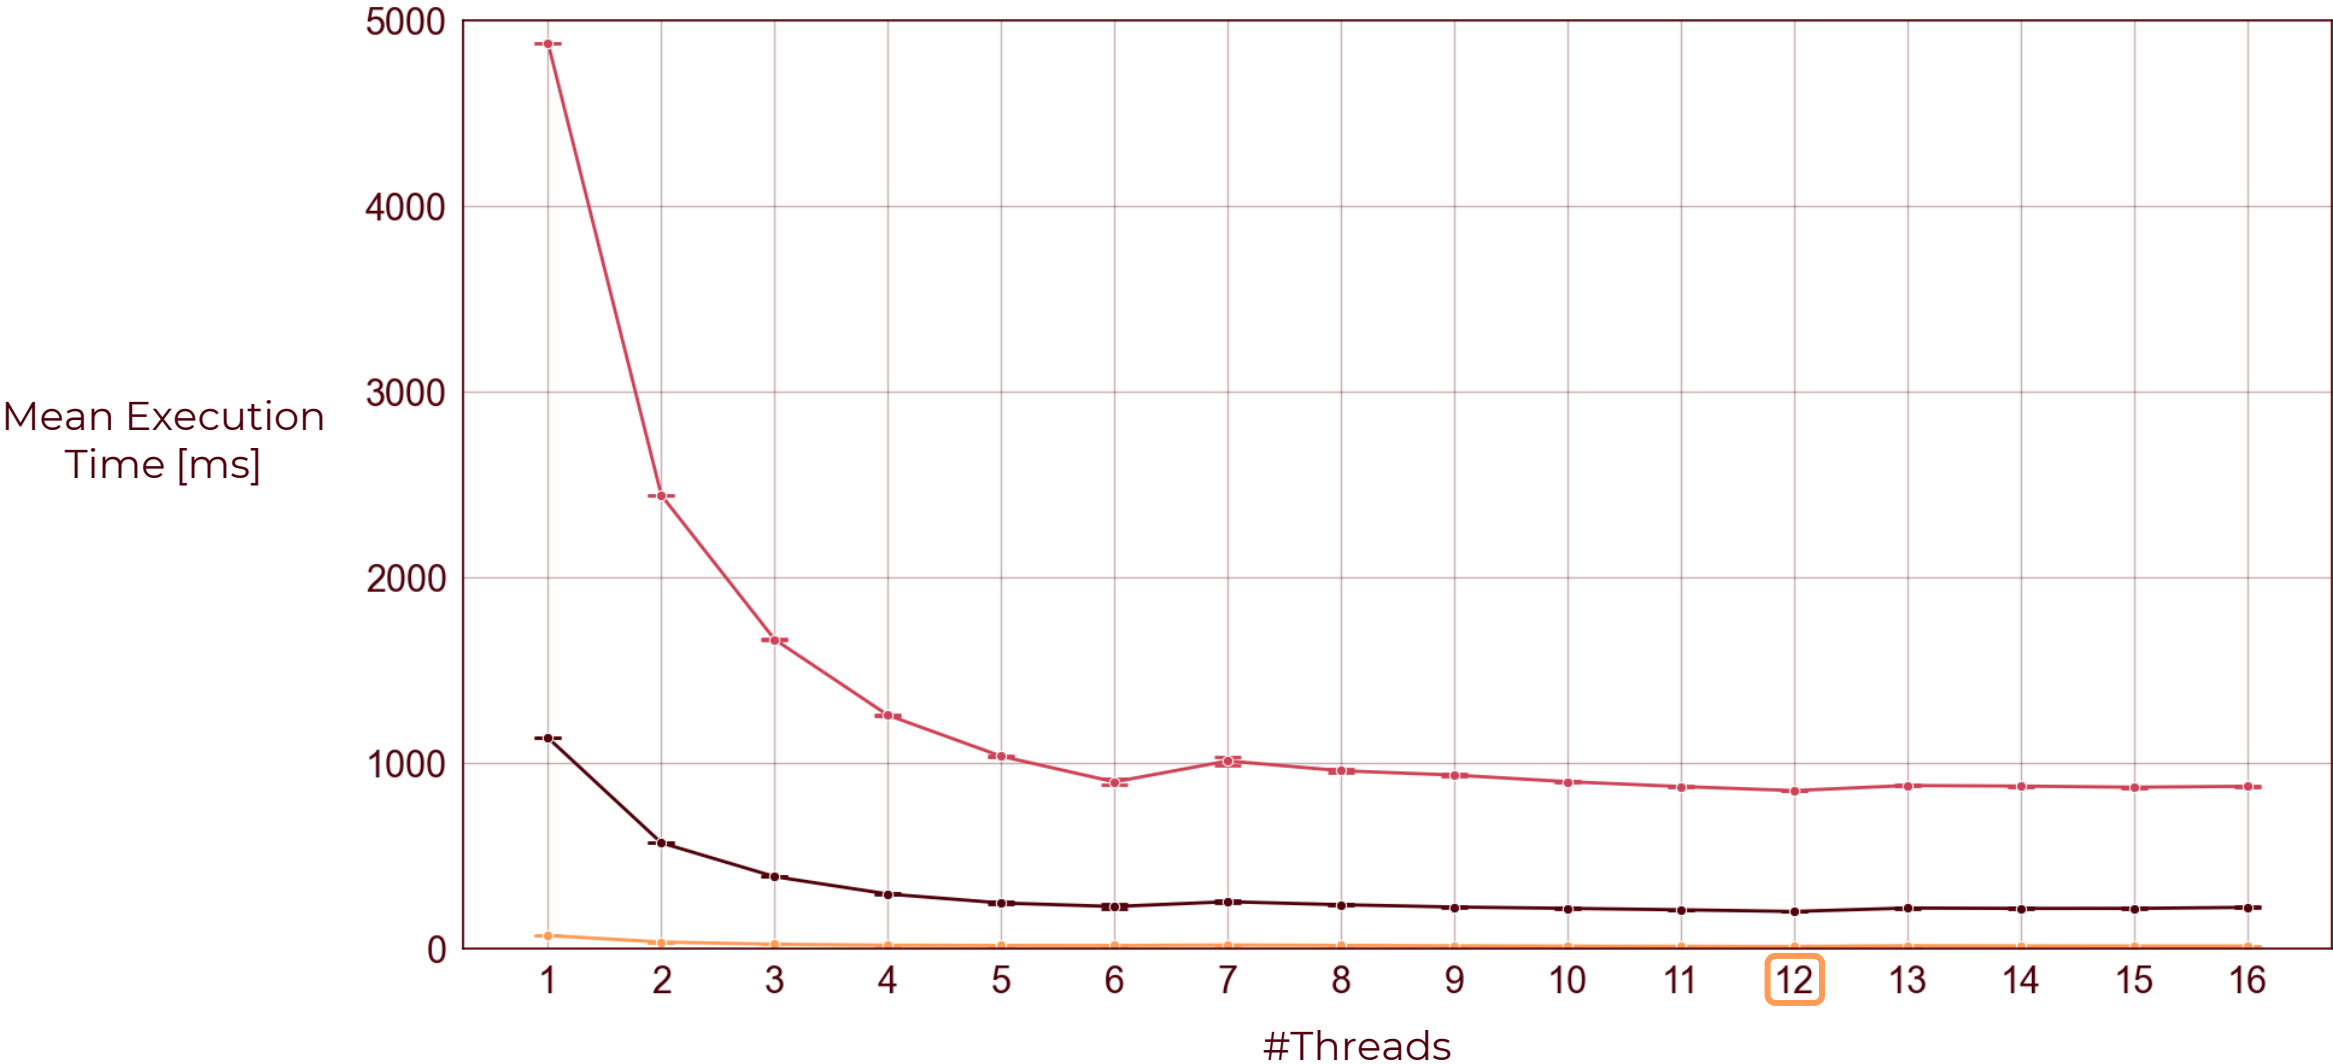
OPTIMIZED VERSION

#ITERATIONS : 30

20 DIGITS 17975734686214396237

18 DIGITS 975734686214396237

16 DIGITS 3934686214396237



SPEED-UP - SCALABILITY


OPTIMIZED VERSION

#ITERATIONS : 30

20 DIGITS 17975734686214396237

18 DIGITS 975734686214396237

16 DIGITS 3934686214396237




(CI omitted due to non-normal distribution)

#Threads

GPU

IMPLEMENTATION






NVIDIA Tesla T4

Streaming Multiprocessors (SM): **40** x
CUDA Cores per SM: **64** =

TOTAL # CUDA Cores: **2560**

Warp Size: **32**



Memory: **16 GB GDDR6 300 GB/s**

Interconnection bandwidth: **32 GB/s**

Max Threads per Block: **1024**

GPU Max Clock Rate: **1590 MHz**

Memory Max Clock Rate: **5001 MHz**

HARDWARE DETAILS



Execution Time

TOOLS

TOTAL
CPU + GPU




`chrono::steady_clock`

GPU
TIME



`cudaEvent_t`

Wall-Clock Time



start



end



NVIDIA **Nsight Compute**




Visual Studio Code

EXECUTION TIME

FIRST VERSION ON GPU

#ITERATIONS : 30
18 DIGITS **975734686214396237**



OPTIMIZATIONS

KERNEL

```
__global__ void findPrimesInRange(...)
```


JUST **1 THREAD** PER WARP WAS WORKING




32 THREADS PER WARP*

*DIVERGENCE CAN HAPPEN BUT ITS EFFECT IS NEGLIGIBLE AS THE AVERAGE NUMBER OF NOT PREDICATED OFF THREADS PER WARP IS **31.64**

32 THREADS x BLOCK



64 THREADS x BLOCK



max_blocks_per_multiprocessor **16**

max_warps_per_multiprocessor **32**


max_warps_per_scheduler **8**



EXECUTION TIME

OPTIMIZED VERSION ON GPU


#ITERATIONS : 30
18 DIGITS **975734686214396237**



EXECUTION TIME

OPTIMIZED VERSION ON GPU

#ITERATIONS : 30
20 DIGITS 17975734686214396237




SPEED-UP

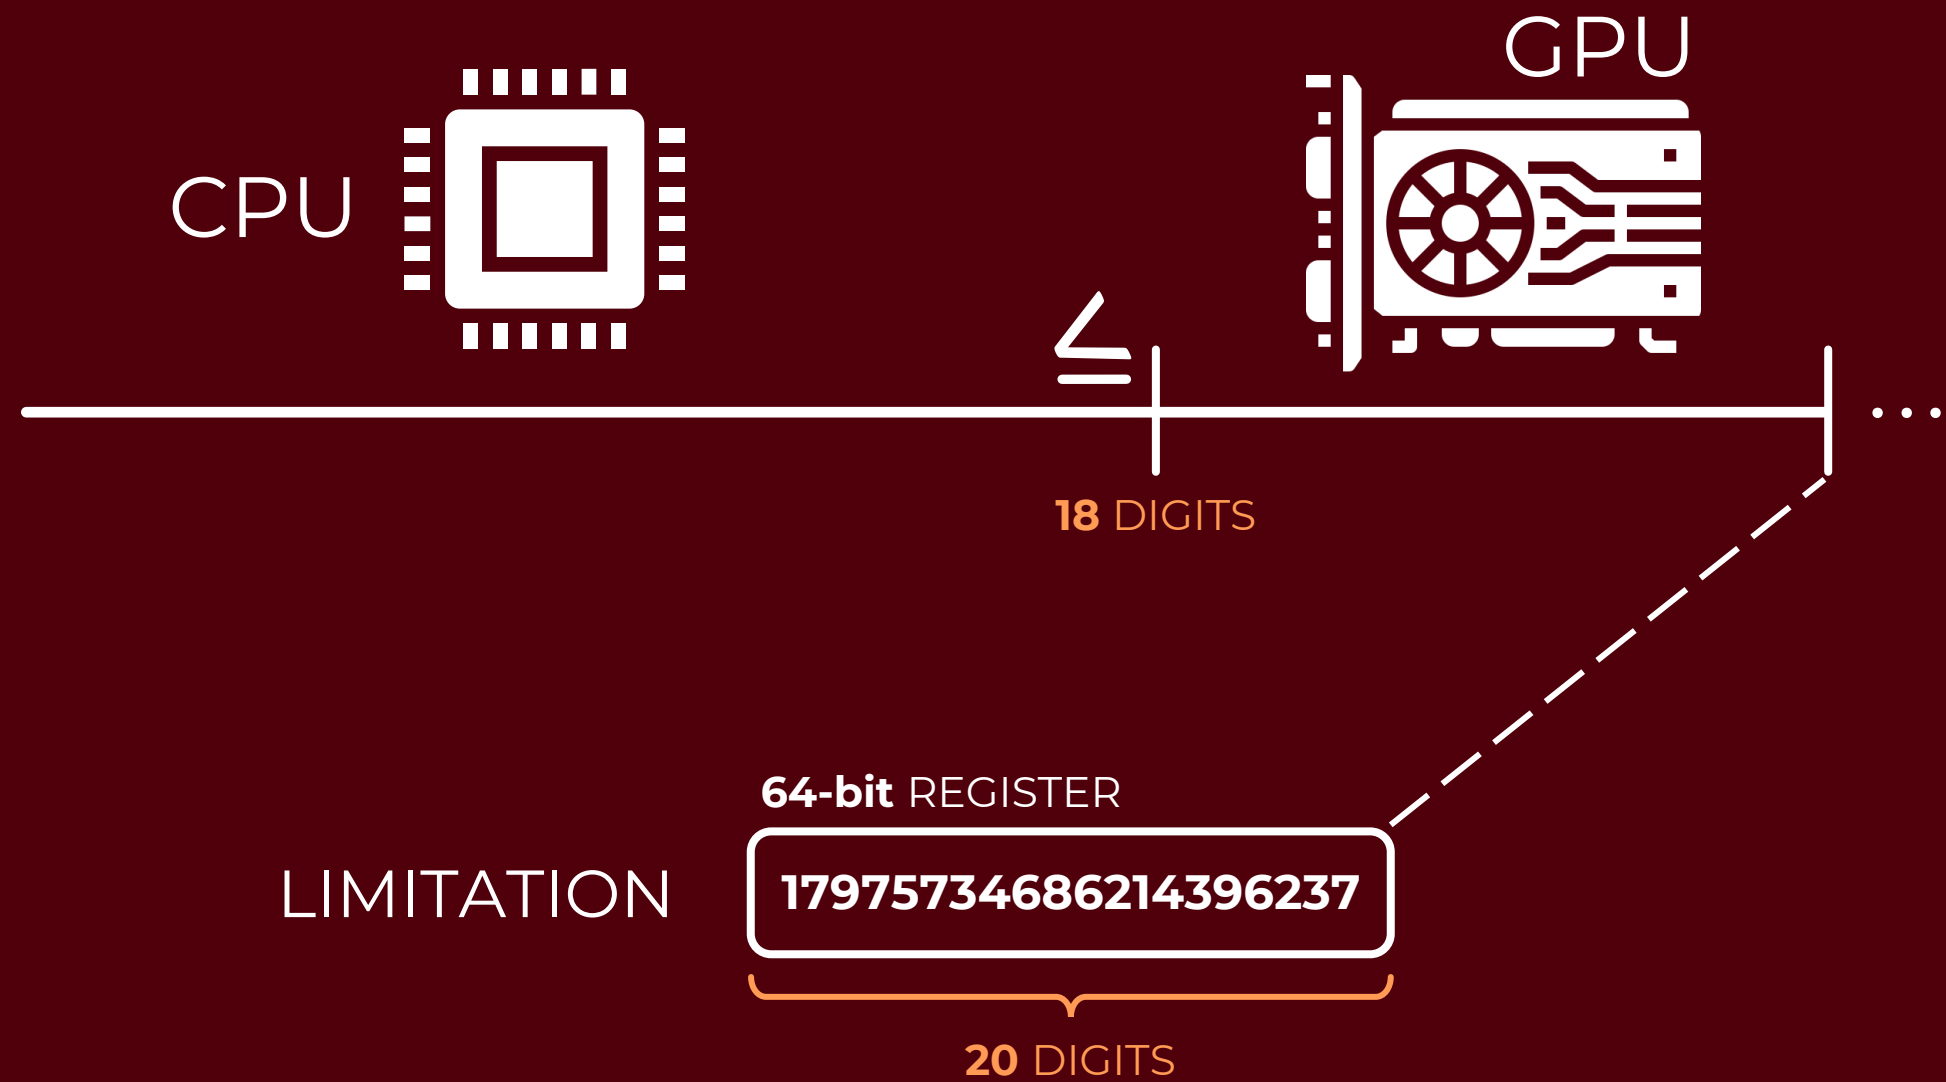
OPTIMIZED VERSION ON GPU


#ITERATIONS : 30

20 DIGITS **17975734686214396237**



CONCLUSIONS





THANKS
FOR THE
ATTENTION

